Programming, Concurrency and Client-Server Computing

School of Computing, Engineering and Mathematics
University of Brighton

Revision Booklet

May, 2012

Contents

Programming Languages

A comparison between Java, Ada, Haskell and Prolog

Concurrency

What it is

Concurrency allows you to effectively utilise available processing power by allowing a computer to perform several different tasks simultaneously (e.g web web server requests, spell-checking whilst typing)

Types of concurrency

Coroutines   With co-operative schedules (or coroutines), processes suspend voluntarily when they need to wait for an event. Processes may also suspend at regular intervals during lengthy processing.

Timeslicing   With pre-emptive scheduling (or timeslicing), processes are suspended as a result of interrupts from external hardware. This can present problems with shared resource access, as it is impossible to predict at what point a process will be suspended.

How it's implemented

Java

Threads   Java's *Thread* class can be used to implement concurrency. Each instance of the class represents a different 'thread' of execution, allowing multiple processes to execute simultaneously.
Useful methods: *Thread.start()*, *Thread.run()* - thread halts on exit from run, *Thread.sleep(time)*, *Thread.interrupt()*, *Thread.get/setPriority()*

Synchronisation   Problems will occur if two threads try to modify/access shared data simultaneously (or rather if one tries to access while one tries to modify, or both try to modify). Using a *synchronized* block prevents data corruption, as only one thread is able to execute it at a time. Each Java object has an internal lock and a queue for waiting threads; if the lock is clear, lock the object and enter the block, if the lock is set, wait. On exit from the block, clear the lock and wake up a waiting thread (if there are any). This isn't very OO, as data is only protected where it is accessed, not where it is defined. The solution is to 'synchronize' an entire method.

Ada

Tasks

Tasks are Ada's equivalent of Java's threads; each separate task representing a separate 'thread' of execution. Tasks can communicate with each other via 'entry' calls. Entry points are declared in the task spec a task declares entries (like procedures) in the task spec in the form of accept statements. An entry call causes a rendezvous between the called task and its caller. The caller waits for the task to accept the call the accept statement waits for a call (see Listing).

Listing 1: Ada's Accept Statement

```
1  loop
2        select
3                accept Show (Message : in String) do
4                        Put_Line(Message);
5                end Show; or
6                accept Stop;
7        exit;
8        end select;
9  end loop;
```

Entries require a client/server relationship, as they only communicate between one task and another. This is inefficient when just accessing shared data (tasks must be rescheduled to actually pass the data).

Protected records

Protected records (added in Ada 95) use shared data much more efficiently. Protected records contain private data which is accessed by public functions, entries and procedures.

Functions are given read-only access to the data and can be executed by multiple tasks simultaneously.

Procedures have read and write access to the data, but can only be executed by a single task at any one time. Additionally, no other tasks can be executing any other functions, procedures or entries while a procedure is executing.

Entries are similar to procedures, but also has a guard condition which suspends the caller until it evaluates to true.

Haskell

Issues Associated With Concurrency

Although concurrency adds the power to do new things, it also brings with it new types of errors.

Deadlock

Deadlock occurs when two (or more) processes require access to an inaccessible certain resource in order to continue. This usually occurs when one process has a lock on some resource which is needed by another process, which itself has a lock on a resource which the first process holds. Neither process are able to relinquish their lock on the problem resources, as they cannot get a lock on the new resources etc. etc.

All four of the following conditions are necessary for deadlock to occur:

1. Tasks need to use a non-shareable resource
   Can be prevented by virtualising resources (e.g. print spooling on disk: printer is non-shareable, disk is shareable), though this is not always possible (e.g. a railway track is not shareable between two trains and cannot be virtualised).

2. Tasks hold onto resources while waiting for extra ones
   Insist that all resources are allocated at once (task cannot proceed until all resources have been granted). This is inefficient as resources will be allocated when not needed.

3. Resources cannot be taken away from tasks by a third party
   See above solution.

4. There is a circular chain of tasks requesting a resource held by another task
   Resources can be prioritised, allocated in priority order. A process must finish using (and release) high priority resources before it can use a lower priority one.

It is not always possible to recover from deadlock. The operating system may check for deadlocks by checking the thread table for circularities. If a deadlock is detected, the OS will kill one of the locked processes until the deadlock is broken. This can obviously have severe implications for the program. Similarly, if deadlocks are not dealt with within the program, it may become unresponsive, forcing the user to kill the program manually.

Livelock

Livelock is similar to deadlock, except that tasks are still able to proceed. However, execution is useless in that tasks will not be able to make any meaningful progress. For example, ethernet, where collisions cause back-offs of exponentially increasing length (this variation in wait time will normally eventually break the lock). Where contention is low, probability of livelock is low enough to ignore (although this is not a suitable response in safety-critical situations). Livelock isn't as easily definable as deadlock, the system may appear to be functioning. Deadlocked threads cannot be scheduled, whereas livelocked ones can. If a 'fair' scheduling algorithm is used, livelock can be avoided (e.g. if a guarantee is made that every request is eventually dealt with).

Starvation

Starvation occurs if one or more tasks are 'starved' of resources by other tasks. This might result from a poor choice of task priorities so that high priority tasks will hog resources

that lower priority tasks also need, meaning that the lower priority tasks are never able to function. Even if the high-priority thread is blocked, the low-priority thread still cant get hold of the resource it needs.

Priority inversion

This occurs if a high-priority task (A) is unable to access a resource which is held by a lower-priority task (B). If A suspends until the resource is free, B is able to proceed. This means that a low-priority process is taking precedence over a high-priority process. If thread A waits in a loop for the resource, the result is a perpetual livelock: B never runs because A is running, but A is always waiting for B.

The Dining Philosophers

> *Five philosophers go to a Chinese restaurant and are seated round a circular table. There is a single chopstick between each pair of philosophers. Each philosopher alternately thinks (which involves putting down any chopsticks that the philosopher is holding) and eats (which involves picking up two chopsticks - only one chopstick can be picked up at a time).*

The dining philosophers is an example of a problem which exhibits the main dangers associated with concurrent systems. **Deadlock:** Each philosopher picks up left chopstick and waits forever for the right chopstick to reappear. **Livelock:** Same as above, but put down left chopstick if right one is unavailable. **Starvation:** Two philosophers can starve another sitting between them if they are never thinking at the same time

Solving the problem: Deadlock can be avoided by providing one extra chopstick or one fewer chair, having one left-handed philosopher, allowing philosophers to snatch chopsticks from each other, philosophers put chopstick down if the other isnt available, or allowing chopsticks to be shared.

Livelock only requires circular waiting for unshareable resources, so can be avoided by picking up left chopstick and put it down again if chopstick not available, or snatching chopstick from neighbours.

Another solution could be to provide a bottle of soy sauce, whereby philosophers can only pick up chopsticks if they have the bottle, sauce is put down sauce after both chopsticks are collected. The bottle ensures mutual exclusion for the critical act of picking up a chopstick.

Distributed Systems

Multi-processors

Becoming increasingly common due to their decreasing cost. Adding extra processors (or cores) is a much more cost-effective way of increasing performance.

Tightly coupled   Multiple processors sharing a single memory (e.g. SMP, quad Pentium motherboards).

Loosely coupled   Distributed networks with no shared memory (e.g. Beowulf).

Majority voting

Problems

Global state

Network issues

Mutual exclusion

Deadlock

How is it implemented

Java

RMI

CORBA

Jini

Networking

Internet fundamentals

TCP/IP

Loosely referred to as 'internet protocols', TCP/IP consists of 5 main layers: hardware, data link (e.g. device driver for Ethernet card), internet (IP), transport (TCP or UDP), application (HTTP etc.). This layered model assumes reliable data transfer at each level; each layer has its own error detection and recovery mechanisms (e.g. by requesting re-transmissions). TCP assumes that reliability is an end-to-end problem: individual packets can be lost or corrupted, and rather than each node passing the message from one endpoint to another, only the transport layer deals with errors detection and recovery. TCP assumes that hosts participate in network issues (routing, error handling, network control). Each destination layer receives the object sent by the same source layer (data link layer receives frames sent by the source data link layer etc.).


Figure 1: The TCP/IP Layered Structure

Ethernet

Ethernet is a broadcast medium which has become the standard for local area networks. Ethernet stations communicate by sending each other data packets (individual blocks of data). Each Ethernet station is given a 48-bit MAC address to identify it on the network. With ethernet, every frame sent is received by every system connected to the Ethernet (including the transmitter).

Frame format: preamble - alternating 0s and 1s for synchronisation (8 octets), destination (6 octets), source (6 octets), frame type (2 octets), data (64-1500 octets), CRC (4 octets)

Carrier Sense Multiple Access with Collision Detection (CSMA/CD): Used to regulate access: A node waiting to transmits monitors the carrier signal waiting for the line to become idle, at which point it starts transmitting. If another node starts transmitting at the same time, the frame is garbled (detectable by the senders). When a collision is detected, each transmitting node backs off for a random period. This randomness reduces chance of another collision (in the case of further collisions, the maximum period for back-off is doubled - exponential). As the network load increases, collisions become more frequent, causing more nodes to back off. This causes the network to slow down, but should not fail completely.

Internet Addresses

Addresses are IPv4 (or 32-bit), consisting of 4 octets expressed as n.n.n.n (e.g. 127.0.0.1). The next-gen IPv6 (128-bit) consists of 6 octets (giving ∼1038 unique addresses).

Address divided into network and machine number: subnet mask specifies how it is divided

ICMP

RARP

BOOTP

DHCP

Transport layer

Name resolution/DNS

How is it implemented

Java

Client/server

Many concurrent systems are implemented as client/server systems, where a server provides a service, and clients communicate with the server to use the service (e.g. web servers and web browsers, time servers providing synchronised time). Java comes with the *java.net* package which contains many network-related classes.

TCP/IP

The Transmission Control Protocol over Internet Protocol allows you to connect to a specific 'port' on a remote machine, and read/write data to the port (like writing to a file). TCP is similar to a phone conversation: you connect, communicate and disconnect.

Sockets

A socket connects to a server, and Input/OutputStreams are used to read and write the data.

Socket servers

Socket servers bind to a particular port and wait for someone to connect, at which point it communicates with the client. Server sockets can be set to timeout (*setSoTimeout(5000);*) so that they don't wait indefinitely for a connection.

UDP (Sockets, Packets)

The User Datagram Protocol doesn't carry the same guarantee for communication: a message isn't guaranteed to reach its destination. Data is sent as an individual 'packet', rather than a continuous stream of data. Having sent a request, a response may be sent (or not). As there isn't necessarily a response, the socket should be set to timeout. UDP is more efficient than TCP, and allows for data to be 'multicasted' to multiple recipients. Multicasting allows broadcasting to recipients whose individual addresses aren't necessarily known. UDP is particularly relevant in video streaming for example, where it doesn't matter if odd packets (or frames) are lost.

Real-time

What is a real-time system

All computer systems model some aspect of the outside world, although the timescale doesn't necessarily match that of the real world. Real time systems are required to conform to timescales imposed by the outside world: they must work as 'fast' (or 'slowly') as the outside world.

Hard real time : very tight deadlines; failure to meet deadlines is catastrophic (e.g. aerospace autopilot).

Soft real time : deadlines are more flexible; failure to meet deadlines is not necessarily catastrophic (e.g. multimedia video display applications, financial payroll systems).

Embedded systems

Real time systems are often computer systems which are part of (embedded in) a larger system e.g. process control, autopilot, manufacturing Embedded often used as a synonym for real time

Characteristics

- Needs to cope with a variety of external events (where software is becoming frequently large and complex)

- Reliable and safe: able to detect and recover from failures

- Need to interact with external hardware

- Need to be able to specify timing requirements (when to perform actions, when to complete actions by, what to do when deadlines are missed, importance of granularity (e.g. IBM PC clock granularity is 55ms = 100yds at 600mph)

- Event-driven rather than process-driven (external events must be dealt with as they occur; event ordering is not generally predictable)

- Generally uses concurrent processes (each event source can be handled by a separate process)

- Must have predictable response times (factors to consider: caching/pipelining affect program speed, worst case is an order of magnitude slower than the average, fast programs & external time)

- Code must be safe (i.e. nothing 'bad' will happen) and live (i.e. something 'good' will eventually happen). 'Eventually' must be able to be upper-bounded, as livelock can particularly be an issue in real-time systems.

Security

Password protection

Encryption

Public-key

Steganography

SSL

Types of attack

Trojan

Virus

Worm

Denial of service

Mail bombing

Phishing

Keylogging

Protection